

CS 598 Practical Statistical Learning

Alexandra Chronopoulou

2026-08-21

Contents

Course Information	5
1 Introduction to Statistical Learning	7
1.1 Examples of Statistical Learning Problems	8
1.2 Supervised Learning Framework	10
1.3 Why Statistical Learning is Challenging	11
1.4 Bias-Variance Trade-Off	12
1.5 Two Toy Examples: k -NN vs. Linear Regression	14
2 Linear Regression Models	21
2.1 Multiple Linear Regression (MLR) Model	22
2.2 MLR Model Fitting	23
2.3 Least-Squares & Normal Equations	27
2.4 Goodness-Of-Fit: R -Square	29
2.5 Linear Transformations on X	30
2.6 Rank deficiency	31
2.7 Hypothesis Testing in MLR	31
2.8 Categorical Variables in MLR	33
2.9 Collinearity	34
2.10 Model Diagnostics	35
2.11 The <code>Birthweight</code> Data Set Example	36

Course Information

Course Description

This course provides an introduction to modern techniques for statistical analysis of complex and massive data. Examples of these include model selection for regression, classification, nonparametric models such as splines and kernel models, regularization, dimension reduction, and clustering analysis. Applications are discussed as well as computation and theoretical foundations.

Course Prerequisites

Knowledge of basic multivariate calculus, statistical inference, and linear algebra. You should be comfortable with the following concepts: probability distribution functions, expectations, conditional distributions, likelihood functions, random samples, estimators, and linear regression models.

Course Learning Outcomes

By the end of the course, you will be able to:

- Use a broad range of methods and techniques in machine learning.
- Have a deeper understanding of major algorithms and techniques in machine learning.
- Build analytics pipelines for regression problems.
- Build analytics pipelines for classification problems.
- Build analytics pipelines for recommendation problems.

Textbook and Readings

There is no required textbook for this course.

Recommended resources for a statistical approach to machine learning are:

- *An Introduction to Statistical Learning with Applications in R* (basic)
- *An Introduction to Statistical Learning with Applications in Python* (basic)

- *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (more advanced)

You may also want to view the Data School YouTube videos associated with the first book as an additional resource.

The detailed syllabus for the course can be found on Coursera.

Chapter 1

Introduction to Statistical Learning

Statistical Learning is a field at the intersection of statistics, computer science, and data science that focuses on developing and understanding models that can *learn* from data. While statistics has long dealt with model building, validation, and prediction, the dramatic growth in the volume, complexity, and dimensionality of data has strained many classical techniques. The introduction of sophisticated algorithms, data-management tools, and optimization methods from computer science has enabled more powerful approaches suited to modern data regimes.

In the supervised setting, learning from data typically means that, given an outcome (quantitative or categorical) and a (usually large) set of features, we construct a **prediction model** (also called a *learner*) that can accurately forecast the outcome for new, previously unseen observations.

We can separate statistical learning problems in two types:


- **Supervised Learning:** when an outcome variable is present to guide the learning process. The two primary examples are:
 - *Regression:* Response is quantitative (continuous or discrete numeric).
 - *Classification:* Response is categorical (binary or multi-class).
- **Unsupervised Learning:** when an outcome variable is not present to guide the learning process. The goal is to discover latent structure or patterns in the data (e.g., clustering, association rules, hidden Markov models, etc.).

1.1 Examples of Statistical Learning Problems

Birthweight Data: A Regression Example

This dataset comes from a study whose goal is to predict the birth weight of a newborn from a collection of baby and parental characteristics (gestation length, length, head circumference, mother's and father's height, mother's weight, parental smoking status, etc.). The study focuses particularly on babies born prematurely (before 40 weeks of gestation).

The scatterplot matrix below shows the pairwise relationships among the variables in the dataset:



Note that three of the predictors (`smoker`, `lowbwt`, and `mage35`) are categorical.

This is a classic example of a supervised regression problem and a natural setting in which a multiple linear regression (MLR) model can be considered. If the MLR model fits well and the key assumptions are reasonably satisfied, it can provide useful predictions of birth weight for new observations.

Trees and Shrubs Data: A Binary Classification Example


In this problem, the goal is to predict whether a woody plant common in the Black Forest region of southwestern Germany is a tree or a shrub (Lederer). A sample of the data is shown below:

Name	Height	Diameter	GrowthRate	Longevity	Type
Silver Fir	55.0	1.50	0.5	500	Tree
European Beech	40.0	1.00	0.5	200	Tree
Common Hazel	5.0	0.10	0.5	70	Shrub
Red Raspberry	1.5	0.01	0.4	10	Shrub
European Spruce	50.0	1.20	0.5	600	Tree

This is a binary classification problem in which we want to use a set of characteristics to predict the class label (tree or shrub).

Handwritten Digits: A Multiclass Classification Example


This is a classic multiclass classification problem in which the goal is to predict the handwritten digit ($0, \dots, 9$) appearing on a postal envelope, based on a 16×16 eight-bit grayscale image (ESL). The practical challenge is to keep the error rate sufficiently low so that mail is not misdirected.



Proteomics Data: A Clustering Example

Proteomics is the large-scale study of proteins. Data of this type are typically generated through laboratory processes such as protein purification or mass spectrometry (a technique that measures the mass-to-charge ratio of ions).

The figure below shows an expression matrix for 437 proteins, of which only 101 are of primary interest (Romanova et al.). A common challenge with such data is the large amount of missing values (shown as grey pixels). The goal is to cluster the proteins into coherent groups for further biological analysis.



These examples illustrate the variety of problems that arise in statistical learning. The birthweight data is a typical supervised regression problem in which we aim to predict a quantitative outcome. The trees-and-shrubs data is a binary classification task, while the handwritten-digit data is a multiclass classification problem. Finally, the proteomics data is an unsupervised learning example whose goal is to uncover latent structure through clustering. Together they show how statistical learning methods can be applied across different types of outcomes and data structures, whether the aim is *prediction* or the *discovery of patterns*.

1.2 Supervised Learning Framework

In this class, our main focus is **supervised learning**. This means that we observe **target** (*outcome, response*) variable Y that we wish to predict using a set of *features* (*predictors*), X . We assume that the relationship between the response Y and the features is given by a function f that is known up to a parameter vector w :

$$\mathbf{Y} = f(\mathbf{X}, w) + \text{varepsilon}$$

When f is fully known except for the value of w (for example when f is linear), then learning problem reduces to **estimating** w . To do so, we collect a **training sample** $\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^n$ and minimize an **objective** function with respect to w :

$$F(w) = \sum_{i=1}^n \text{loss}(y_i, f(\mathbf{x}_i, w))$$

The **loss** function measures the discrepancy between the model prediction $f(\mathbf{x}_i, w)$ and the observed outcome $\mathbf{y}_i$. The choice of loss is made by the practitioner; the most common examples are:

- **Squared Error Loss**

$$\text{loss} = \left[y_i - f(x_i; w) \right]^2$$

- **Absolute Error Loss**

$$\text{loss} = \left| y_i - f(x_i; w) \right|$$

- **Log Loss** (binary classification)

$$\text{loss} = -y_i \log f(x_i; w) - (1 - y_i) \log(1 - f(x_i; w))$$

- **0-1 Loss**

$$\text{loss} = \mathbf{1}_{\{y_i \neq \hat{y}_i\}} \quad \text{where} \quad \hat{y}_i = f(x_i; w)$$

Depending on the chosen loss, the minimizer of $F(w)$ may or may not admit a closed-form expression. When a closed form is unavailable we resort to numerical optimization algorithms. In many cases gradient-based methods (in particular gradient descent and its variants) can be used to obtain a good solution.

1.3 Why Statistical Learning is Challenging

In the framework of the previous section we reduced the learning problem to the estimation of a parameter vector w , under the assumption that the true function f is known up to this parameter. In practice, however, we rarely know the exact form of f . We must therefore *approximate* it by some function $\hat{f}$. This approximation introduces an *additional source of error*.

Moreover, the objective function we minimize is computed on a *training* sample. While this quantifies how well the model fits the observed data, the real goal of statistical learning is good performance on **new, unseen** data. In other words, we aim to minimize the **test error** (also called generalization error), not the training error.

Statistical learning is therefore challenging for several reasons:

- The training error typically underestimates the test/generalization error.
- A model may perform well on the training data yet poorly on future data because of overfitting. The number of parameters required to approximate f can be large—sometimes larger than the number of available observations.
- As the number of model parameters grows, the test error often increases substantially.

Two simple illustrations make these points concrete:

1. **Classification: 1-Nearest Neighbour**

The one-nearest-neighbour classifier achieves zero training error (it predicts every training point perfectly) but typically performs poorly on test data. An instructive demonstration of how increasing dimensionality affects the performance of linear classifiers can be found here.

2. **Regression: the p versus n problem**

Consider the linear model

$$\mathbf{y}_{n \times 1} = \mathbf{X}_{n \times p} \mathbf{w}_{p \times 1} + \varepsilon,$$

where $\mathbf{y}$ is the response vector, $\mathbf{X}$ is the design matrix with p predictors, and $\mathbf{w}$ contains the coefficients to be estimated.

- When $p < n$ we have more equations than parameters; classical regression methods usually work well.
- When $p = n$ we have exactly as many equations as parameters and can achieve a perfect fit on the training data.
- When $p > n$ we have fewer equations than parameters. In this high-dimensional regime classical methods break down and specialized techniques are required.


In summary, statistical learning is challenging because we must approximate an unknown function f , control the gap between training and test error, and often operate in regimes where the number of parameters is large relative to the sample size. Successfully addressing these issues is essential for building models that generalize well to new data.

1.4 Bias-Variance Trade-Off

We have already noted that *two sources of error* affect the quality of our predictions. In statistical learning we must balance:


- **Bias:** the systematic difference between the average of our estimated function $\hat{f}$ and the true underlying function f . High bias typically leads to under-fitting and an inaccurate model.
- **Variance:** the amount by which $\hat{f}$ would change if we estimated it using a different training sample. High variance typically leads to over-fitting and an unreliable model.

A helpful analogy is to think of the learning problem as a **target** whose *center represents the true model*.



How close the fitted values lie to the center is measured by bias. If the points systematically miss the center, the model is a poor approximation of f . Increasing the sample size will not correct a large bias. The spread of the fitted values around their average is measured by variance. When the points are tightly clustered the variance is low; when they are widely scattered the variance is high.

Ideally we would like to reduce both bias and variance simultaneously. In practice this is not possible, so we must trade one against the other. Remembering that our ultimate goal is to minimize the generalization (test) error, not the training error, we observe the following typical behavior:



- When the model is highly complex (many parameters), the training error tends to decrease. However, the model often over-fits the training data and therefore exhibits high variance on new data, resulting in poor generalization.

- When the model is very simple (few parameters), it tends to under-fit the data. This produces large bias and, again, poor generalization performance.

In this class we will discuss:

- (i) flexible modeling techniques that help reduce bias, and
- (ii) practical strategies for achieving a good bias-variance trade-off.

Two important approaches that we will study in detail are:

- **Regularization:** constrain the parameters to a lower-dimensional space that is determined adaptively by the data.
- **Ensemble:** average many low-bias, high-variance models; the averaging step reduces variance while largely preserving the low bias.

1.5 Two Toy Examples: k -NN vs. Linear Regression

Before we conclude this introductory chapter, we examine two simple supervised learning methods:

- (i) k -Nearest Neighbors (k -NN)
- (ii) Linear regression

We will compare their behavior on simple examples and use them to illustrate the bias-variance trade-off.

1.5.1 k -Nearest Neighbors

In the k -nearest neighbors (k -NN) method we predict the response at a point $\mathbf{x}$ by using the k training observations that are closest to $\mathbf{x}$. The k -NN fit is

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{x_i \in N_k(\mathbf{x})} y_i$$

where $N_k(\mathbf{x})$ denotes the neighborhood of $\mathbf{x}$ consisting of the k nearest training points.

- In **regression**, $\hat{y}(\mathbf{x})$ is a local average of the neighboring responses.

- In **classification**, $\hat{y}(\mathbf{x})$ is the majority class (or the class probability) within the neighborhood.

Two practical challenges that arise are choosing the neighborhood size k and choosing the distance metric that defines the neighborhood.

The value of k directly controls model complexity, which is roughly proportional to n/k .

- When $k = 1$ the prediction at each training point x_i is exactly y_i , yielding zero training error.
- When $k = n$ every neighborhood contains the entire training sample, so the prediction is constant for all $\mathbf{x}$.

The “default” metric is Euclidean distance

$$d(\mathbf{x}, \tilde{\mathbf{x}}) = \sum_{j=1}^p w_j (x_j - \tilde{x}_j)^2,$$

where the weights w_j may be learned from the data.

1.5.2 Linear Regression

Given an input vector $\mathbf{x}^\top = (x_1, \dots, x_p)$, we approximate the response by a linear function

$$f(\mathbf{x}) \approx \beta_0 + \sum_{j=1}^p x_j \beta_j$$

The coefficients β_j are estimated by ordinary least squares (OLS), i.e., by minimizing the residual sum of squares (*objective function*)

$$\min_{\beta_0, \dots, \beta_p} \sum_{i=1}^n \left(y_i - \beta_0 - x_{i1}\beta_1 - \dots - x_{ip}\beta_p \right)^2$$

Under standard conditions the solution has a closed form, and the fitted value at the i -th observation is

$$\hat{y}_i = \hat{y}(x_i) = x_i^T \hat{\beta}$$

We postpone the algebraic details until Week 2.

Linear Regression in a Classification Setting

Linear regression can also be applied to binary classification problems in which $Y \in \{0, 1\}$. One simply predicts `class 1` when the fitted value $\hat{f}(\mathbf{x})$ exceeds 0.5 and `class 0` otherwise.

This approach has well-known drawbacks: the squared-error loss is not a natural measure for binary outcomes, and a linear function can produce fitted values outside the interval $[0, 1]$. For these reasons logistic regression is the preferred linear method for classification. It models the log-odds

$$\log \frac{p(\mathbf{x})}{1 - p(\mathbf{x})} \approx \beta_0 + \sum_{j=1}^p x_j \beta_j$$

We will return to logistic regression in Week 8.

1.5.3 A Simulated (Binary) Classification Example

Consider a response variable G that takes two values (0 = BLUE or 1 = ORANGE). We generate 200 observations (100 in each class) and compare linear regression and k -nearest neighbors **as classifiers**.


The code used to produce the simulation and all figures is available in R and Python.

Linear Regression Classifier

We first fit a linear regression model, treating the binary response as if it were continuous. The continuous fitted values $\hat{Y}$ are then converted to class predictions $\hat{G}$ by the simple rule

$$\hat{G} = \begin{cases} \text{Blue,} & \text{if } \hat{Y} > 0.5 \\ \text{Orange,} & \text{otherwise} \end{cases}$$

Because the problem is two-dimensional, the decision boundary is a *straight line*.




The fitted boundary (black line in the figure below) is the set of points satisfying $\mathbf{x}^\top \hat{\beta} = 0.5$. In the simulation the true blue region lies above this line and the orange region below it. Many points on both sides of the boundary are nevertheless misclassified. The linear decision boundary is smooth but rigid; it cannot adapt to local structure in the data.

k -Nearest-Neighbor Classifier


We now apply k -nearest neighbors to the same data. Using $k = 15$, $\hat{Y}(\mathbf{x})$ is the proportion of blue observations among the 15 nearest neighbors of $\mathbf{x}$. Class labels are again assigned by the 0.5 threshold:

$$\hat{G} = \begin{cases} \text{Blue, if } \hat{Y} > 0.5 \\ \text{Orange, otherwise} \end{cases}$$

(The prediction is therefore a majority vote among the 15 neighbors.)



The resulting decision boundary is far more irregular and follows local clusters of blue and orange points. Consequently the number of misclassifications on the training data is smaller than with linear regression. If we take the extreme case $k = 1$, each point is classified by its single nearest neighbor:




The boundary becomes extremely rough and the training misclassification rate drops nearly to zero. Whether this is desirable, however, can be judged only by examining the *generalization error*.

Training versus Test Error

Up to this point we have evaluated both methods on the same data that were used to fit them. A method such as 1-NN therefore appears to have zero error. Fair comparison requires an independent test set. The figure below shows misclassification error on a test set of 10,000 observations as a function of model degrees of freedom. Several k -NN fits (varying k) are compared with the linear

regression fit.



The magenta curve is the test error and the blue curve is the training error of k -NN. A 5-fold cross-validation procedure selected the optimal k (marked by a diamond). Linear regression results appear as the magenta and blue triangles at 3 degrees of freedom (the dimension of the fitted linear model). Linear regression has only a few parameters and therefore relatively low variance, but the linearity assumption is restrictive for this problem and produces high bias.

In contrast, k -NN makes almost no assumption about the shape of the decision boundary (apart from a mild local-smoothness condition). This flexibility yields low bias but high variance; the extreme case $k = 1$ corresponds to roughly n parameters and severely overfits. It can be shown that k -NN is consistent when $k, n \rightarrow \infty$ with $k/n \rightarrow 0$. The price of that consistency is the high variance that appears when the effective number of parameters becomes large.

Chapter 2

Linear Regression Models

A regression model is a model for the **conditional expectation** of the response/outcome variable given a set of predictors/features. A linear regression model specifically assumes that $E(Y | \mathbf{X} = \mathbf{x})$ is a linear function of the inputs $X_1, X_2, \dots, X_p$, i.e.,


$$E(Y | \mathbf{X} = \mathbf{x}) \sim \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$

To describe the properties and characteristics of the multiple linear regression (MLR) model, we introduce the **Birthweight** example.

Introduction to the Birthweight Example

The goal of the birthweight study is to **predict** the birthweight (*response/outcome*) of a newborn baby given a set of *predictors/features*, specifically for babies born prematurely.

The features include baby-related characteristics, such as length, weight, head circumference, gestation period, and parent characteristics, such as mother's/father's height, smoking habits, mother's pre-pregnancy weight, mother's/father's age, etc.



The **predictors** we consider may be:

- quantitative, such as length, weight, or gestation period;
- qualitative/categorical, such as smoking habits (yes/no); or
- transformations of quantitative inputs. Examples include power transformations (e.g., X_3^3), continuous functions of variables (e.g., $\log X_5$), polynomial expansions (e.g., $\sum_{k=1}^n c_k X_7^k$), and interactions (e.g., $X_1 \cdot X_2$).

2.1 Multiple Linear Regression (MLR) Model

The model formulation is given by:

Multiple Linear Regression Model

$$Y = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p + \varepsilon$$

where

- β_0 is the intercept
- β_j is the regression coefficient associated with predictor X_j
- ε is the error term. The usual assumptions for the error terms are $\varepsilon \sim IID(0, \sigma^2 \mathbf{I})$

Given the **training data** $\{x_{i1}, x_{i2}, \dots, x_{ip}; y_i\}_{i=1}^n$, we can rewrite the model as

$$y_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \varepsilon_i, \quad i = 1, \dots, n$$

where n is the sample size. Using the following matrix representation for the

response and the features:

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{pmatrix} = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{pmatrix} \quad \mathbf{X} = \begin{pmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & x_{21} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{pmatrix}$$

the model equation can be written in as

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & x_{21} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{pmatrix} + \begin{pmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{pmatrix}$$

which leads to the concise form

$$\begin{array}{ccccccc} \mathbf{y}_{n \times 1} & = & \mathbf{X}_{n \times (p+1)} & \begin{matrix} (p+1) \times 1 \\ \uparrow \\ \text{Coefficients} \end{matrix} & + & \begin{matrix} n \times 1 \\ \uparrow \\ \text{Error Term} \end{matrix} \\ \uparrow & & \uparrow & & & \uparrow \\ \text{Response} & & \text{Design Matrix} & & & \end{array}$$

where n is the sample size, and $p+1$ is number of columns of $\mathbf{X}$ (the p predictors plus the intercept).

2.2 MLR Model Fitting

Given the training data $\{x_{i1}, x_{i2}, \dots, x_{ip}; y_i\}_{i=1}^n$, we want to estimate β , i.e. express:


$$\hat{\beta} = (\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p)^T$$

as a **function of the data**.

The estimated regression coefficients $\hat{\beta}$ are obtained by minimizing the *Residual Sum of Squares* (RSS):

$$RSS = \|\mathbf{y} - \mathbf{X}\beta\|^2 = (\mathbf{y} - \mathbf{X}\beta)^T(\mathbf{y} - \mathbf{X}\beta)$$

The RSS minimizes the (Euclidean) distance of the points from the regression surface:



This approach makes minimal assumptions, since it only requires that the pairs (x_i, y_i) are *random draws from their population*. Specifically, it makes **no assumptions** about the underlying distribution of the response. It simply finds the **best linear fit**, without assuming that the model is valid. In fact, the true underlying model does not even need to be linear for this method to produce estimators. Therefore, it often works well regardless of how the data were obtained. If a linear model provides a good approximation to a non-linear truth, then the regression surface can also be viewed as a measure **lack-of-fit**.

Least-Squares Method

We want to estimate the **vector** of coefficients by minimizing the *sum of squared residuals*:

$$RSS = \|y - \mathbf{X}\beta\|^2 = (y - \mathbf{X}\beta)^T(y - \mathbf{X}\beta)$$

We take derivatives with respect to β and set them equal to zero:

$$\begin{aligned} \frac{\partial RSS}{\partial \beta} &= \mathbf{0}_{p \times 1} \Leftrightarrow \\ -2 \mathbf{X}_{p \times n}^T (y - \mathbf{X}\beta)_{n \times 1} &= \mathbf{0}_{p \times 1} \end{aligned}$$

This leads to the so-called **Normal Equations**

$$\mathbf{X}^T(y - \mathbf{X}\beta) = \mathbf{0}$$

Solving the Normal Equations

$$(\mathbf{X}^T \mathbf{X}) \beta = \mathbf{X}^T y$$

leads to the

Least Squares Estimators in MLR

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y$$

We assume that **the rank of $\mathbf{X}$ is $p + 1$** , i.e. no columns of $\mathbf{X}$ are a linear combinations of the other columns. *Since $\mathbf{X}$ has full column rank, the inverse of $(\mathbf{X}^T \mathbf{X})$ exists.*

Fitted, Predicted Values & Residuals

The **fitted value** of y_i at $x_i = (x_{i1}, x_{i2}, \dots, x_{ip})$ is computed as:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \hat{\beta}_2 x_{i2} + \dots + \hat{\beta}_p x_{ip}$$

More generally, using matrix formulation, the **fitted values** of $\mathbf{y}$ are given by

$$\begin{aligned} \hat{\mathbf{y}}_{n \times 1} &= \mathbf{X} \hat{\beta} \\ &= \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y \\ &= \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y := \mathbf{H}_{n \times n} y_{n \times 1} \end{aligned}$$

where we define the *hat matrix*

The Hat Matrix

$$\mathbf{H}_{n \times n} = \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

It is called the **hat matrix** because it produces the “*y-hat*” values: $\hat{y} = \mathbf{H}t$.

The **predicted value** of y at a new point $x_i^* = (x_{i1}^*, x_{i2}^*, \dots, x_{ip}^*)$ is

$$\hat{y}^* = \hat{\beta}_0 + \hat{\beta}_1 x_{i1}^* + \hat{\beta}_2 x_{i2}^* + \dots + \hat{\beta}_p x_{ip}^*$$

In matrix form, the **predicted values** are

$$\hat{\mathbf{y}}^* = \mathbf{X}^* \hat{\beta}$$

Note: The key difference between the fitted and predicted values is that the vector $\mathbf{x}$ belongs to the training data (it has already been observed), whereas $\mathbf{x}^*$ is an **unobserved** feature vector that is **independent** of the training data.

The **residual** at observation i is

$$\begin{aligned} r_i &= y_i - \hat{y}_i \\ &= y_i - \hat{\beta}_0 - \hat{\beta}_1 x_{i1} - \hat{\beta}_2 x_{i2} - \dots - \hat{\beta}_p x_{ip} \end{aligned}$$

In matrix form the residual vector is

$$\begin{aligned}\mathbf{r}_{n \times 1} &= \mathbf{y} - \hat{\mathbf{y}} \\ &= \mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}} = \mathbf{y} - \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{y} - \mathbf{H}\mathbf{y} = (\mathbf{I} - \mathbf{H})\mathbf{y}\end{aligned}$$

The residuals $\mathbf{r}$ are used to estimate the **error variance**:

$$\hat{\sigma}^2 = \frac{1}{n - p - 1} \sum_i r_i^2 = \frac{RSS}{n - p - 1}$$

The denominator $n - p - 1$ is the *degrees of freedom of the residuals*.

Properties of the Residuals

The least-squares estimator $\hat{\boldsymbol{\beta}}$ satisfies the **normal equations**

$$\mathbf{X}^T(\mathbf{y} - \hat{\mathbf{y}}) = \mathbf{X}^T(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) = \mathbf{0}$$

From this we obtain the following properties of the residual vector $\mathbf{r}$:

- The residual vector is orthogonal to every column of $\mathbf{X}$:

$$\begin{aligned}\mathbf{X}^T \mathbf{r} &= \mathbf{X}^T(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) \\ &= \mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \hat{\boldsymbol{\beta}} \\ &= \mathbf{X}^T \mathbf{y} - (\mathbf{X}^T \mathbf{X})(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{0}\end{aligned}$$

- The residual vector is also orthogonal to the vector of fitted values:

$$\hat{\mathbf{y}}^T \mathbf{r} = \hat{\boldsymbol{\beta}}^T \mathbf{X}^T \mathbf{r} = \mathbf{0}$$

Thus, the residual vector $\mathbf{r}$ is **orthogonal to each column of $\mathbf{X}$ and to $\hat{\mathbf{y}}$** .

When the model is used for prediction at a new point $\mathbf{x}^*$, two different interval estimates are commonly reported: (i) a confidence interval for the mean response $E(Y | \mathbf{x}^*)$, and (ii) a prediction interval for a new individual observation Y^* . The prediction interval is wider because it must also account for the irreducible error variance σ^2 .

2.3 Least-Squares & Normal Equations

In this section, we focus on the system of linear equations that we solve to obtain the least-squares estimators for β . Recall, that we want to find a vector $\hat{\beta}$ that minimizes

$$\min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2$$

Any vector that provides a minimum value of this expression is called a *least-squares solution*.

The set of all least squares solutions is precisely the **set of solutions** to the normal equations

$$(\mathbf{X}^T \mathbf{X})\beta = \mathbf{X}^T \mathbf{y}$$

There is a *unique* solution **if and only if** $\text{rank}(\mathbf{X}) = p + 1$ in which case $(\mathbf{X}^T \mathbf{X})$ is invertible.

System of Linear Equations

Linear Algebra Review Let us review a few facts from linear algebra related to solutions of a system of linear equations. For simplicity of notation, we use the *generic* system of equations:

$$\mathbf{A}z = c$$

where

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1k} \\ a_{21} & a_{22} & \cdots & a_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mk} \end{pmatrix}, \quad z = \begin{pmatrix} z_1 \\ z_2 \\ \vdots \\ z_m \end{pmatrix}, \quad c = \begin{pmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{pmatrix}$$


Denote by $\mathbf{A}_i = (a_{1i}, a_{2i}, \dots, a_{mi})^T$ the vector representing the *i*th *column* of $\mathbf{A}$. Then, $\mathcal{C}(\mathbf{A})$ is the space generated by the columns of $\mathbf{A}$, i.e. the span of $\{\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_k\}$.

Definition

The **span** of a collection of vectors $(\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_k)$ is the set of all linear combinations of these vectors:

$$\begin{aligned} & \text{span}(\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_k) \\ &= \left\{ d_1 \mathbf{A}_1 + d_2 \mathbf{A}_2 + \dots + d_k \mathbf{A}_k, \text{ for any constants } d_1, \dots, d_k \in \mathbb{R} \right\} \end{aligned}$$

The column space of $\mathbf{A}$, $\mathcal{C}(\mathbf{A})$, is a **subspace** in $\mathbb{R}^m$. A linear subspace of $\mathbb{R}^m$ can be thought of as a “flat” surface within $\mathbb{R}^m$; it is a collection of vectors that is closed under linear combinations. Illustrations of subspaces of $\mathbb{R}^n$ are shown below:



Solving the System of Equations $\mathbf{A}z = c$

For the system $\mathbf{A}z = c$ to have a solution, c must lie in the column space of $\mathbf{A}$, i.e. $c \in \mathcal{C}(\mathbf{A})$. This follows directly from the definition of matrix multiplication:

$$\mathbf{A}z = c \Leftrightarrow c = z_1\mathbf{A}_1 + \dots + z_k\mathbf{A}_k$$

When $c \notin \mathcal{C}(\mathbf{A})$, we instead seek the vector $\hat{c} \in \mathcal{C}(\mathbf{A})$ that is closest to c . In that case, the equation $\mathbf{A}z = \hat{c}$ has a solution, and $\hat{c}$ comes as close as possible to the original vector c . To obtain $\hat{c}$ we **project c orthogonally** onto $\mathcal{C}(\mathbf{A})$ by multiplying both sides by $\mathbf{A}^T$:

$$\mathbf{A}^T\mathbf{A}z = \mathbf{A}^Tc$$

* These are the normal equations.

In the sketch to the right, $\mathcal{C}(\mathbf{A})$, the space spanned by the columns of $\mathbf{A}$, is a *flat* surface, and c is a point lying off that flat surface. The **shortest distance** from the point c to the plane $\mathcal{C}(\mathbf{A})$ is the line segment *orthogonal* to the plane.

The *normal equations* locate the point in $\mathcal{C}(\mathbf{A})$ that is closest to c by means of an orthogonal projection.

Geometric Representation of LS

Translating the preceding discussion into the linear-regression setting:

- $\mathcal{C}(\mathbf{X})$ is the subspace *spanned* by the columns of the design matrix, i.e. by the predictors $X_1, X_2, \dots, X_p$. It is a *flat* subspace of $\mathbb{R}^n$.

- The response $\mathbf{y}$ is a vector lying **off** this subspace, and the fitted value $\hat{\mathbf{y}}$ is the orthogonal projection of $\mathbf{y}$ onto $\mathcal{C}(\mathbf{X})$.
- The residual vector $\mathbf{r} = \mathbf{y} - \hat{\mathbf{y}}$ is **orthogonal** both to $\hat{\mathbf{y}}$ and to the entire subspace $\mathcal{C}(\mathbf{X})$.

In other words, the least-squares method decomposes the data vector $\mathbf{y}$ into two orthogonal components:

$$\mathbf{y}_{n \times 1} = \hat{\mathbf{y}}_{n \times 1} + \mathbf{r}_{n \times 1}$$

2.4 Goodness-Of-Fit: R -Square

A measure of how well the model fits the data is the **R -square** (also called **coefficient of determination** or the *percentage of variance explained*). It is defined as

$$\begin{aligned} R^2 &= \frac{\sum_i (\hat{y}_i - \bar{y})^2}{\sum_i (y_i - \bar{y})^2} = \frac{\text{distance of model from grand mean}}{\text{distance of observations from grand mean}} \\ &= \frac{\|\hat{\mathbf{y}} - \bar{\mathbf{y}}\|^2}{\|\mathbf{y} - \bar{\mathbf{y}}\|^2} \\ &= \frac{\|\mathbf{y} - \bar{\mathbf{y}}\|^2 - \|\hat{\mathbf{y}} - \mathbf{y}\|^2}{\|\mathbf{y} - \bar{\mathbf{y}}\|^2} \\ &= 1 - \frac{\|\hat{\mathbf{y}} - \mathbf{y}\|^2}{\|\mathbf{y} - \bar{\mathbf{y}}\|^2} := 1 - \frac{RSS}{TSS} \end{aligned}$$

The second-to-last equality follows from the orthogonal decomposition


$$\underbrace{\|\mathbf{y} - \bar{\mathbf{y}}\|^2}_{\text{total variation}} = \underbrace{\|\hat{\mathbf{y}} - \bar{\mathbf{y}}\|^2}_{\text{model variation from mean}} + \underbrace{\|\hat{\mathbf{y}} - \mathbf{y}\|^2}_{\text{error}}$$

Properties of R^2

- $0 \leq R^2 \leq 1$
- R^2 is invariant of any location or scale change of Y or of the predictors X .
- R^2 alone does not tell us much about the quality of the LS fit.
- A small R^2 does not necessarily imply that the LS model is poor.

- Adding a new predictor, even if it is randomly generated and completely unrelated to Y , decreases the RSS and therefore increases R^2 .

These relationships are illustrated in the two plots below:



2.5 Linear Transformations on X

Suppose we have a linear regression model of Y on the predictors X . If we **scale** or **shift** a predictor, the fitted values $\hat{y}$ and R^2 remain unchanged, but the LS estimators $\hat{\beta}$ change.

Some Examples

- Scaling a predictor, i.e. $\tilde{x}_{i2} = 2x_{i2}$ or $\tilde{x}_{i3} = x_{i3}/4$, does not affect $\hat{y}$ or R^2 .
- In the **birthweight** example, the length of a baby is recorded in cm. Converting the length column to inches by the transformation $\tilde{X}_1 = X_1/2.54$ leaves the predicted values unchanged.

These statements remain true, if we apply any linear transformation of the p predictors **as long as the transformation does not change the rank of X** .

2.6 Rank deficiency

The design matrix $\mathbf{X}$ is an $n \times (p + 1)$ matrix. If this matrix is **not of full rank** (i.e., its columns are not linearly independent), the matrix $\mathbf{X}^T \mathbf{X}$ cannot be inverted. This means that there is redundancy among the columns of $\mathbf{X}$: at least one column can be written as a linear combination of other columns. Consequently, $\mathbf{X}^T \mathbf{X}$ is *singular* and the LS solution is not unique (an *identifiability problem*).

Nevertheless, the column space $\mathcal{C}(\mathbf{X})$ remains well-defined, so the fitted vector $\hat{y}$ is still well-defined and can be computed. **R** and **Python** handle this situation, and the resulting model can still be used for prediction.

2.7 Hypothesis Testing in MLR

When fitting a regression model, we often want to understand which predictors are statistically significant, or which of two models (a larger one or a smaller one) is more appropriate –either for prediction or for estimation.

Testing for the statistical significance of one or more predictors can be formulated either as a test that the corresponding coefficient(s) equal zero, or as a model-comparison test. In this course we summarize all such testing questions with the following hypothesis test:

$$\begin{cases} H_0 : \text{Reduced Model with } p_0 \text{ coefficients} \\ H_\alpha : \text{Full (larger) Model with } p_\alpha \text{ coefficients} \end{cases}$$

Obviously, $p_\alpha > p_0$, since the reduced model is a special case of the full model. Consequently we expect the reduced model to have a larger residual sum of squares: $RSS_0 \geq RSS_\alpha$. The main tool for this comparison is the **partial F test**:

Partial F test

Hypothesis Test:

$$\begin{cases} H_0 : \text{Reduced Model with } p_0 \text{ coefficients} \\ H_\alpha : \text{Full/ Larger Model with } p_\alpha \text{ coefficients} \end{cases}$$

Test Statistic:

$$F = \frac{(RSS_0 - RSS_\alpha)/(p_\alpha - p_0)}{RSS_\alpha/(n - p_\alpha)} \sim F_{p_\alpha - p_0, n - p_\alpha}$$

Decision Rule: Reject H_0 , if the F test statistic is large (or larger than $F_{p_\alpha - p_0, n - p_\alpha}$), i.e. the variation missed by the reduced model, when being compared with the error variance, is significantly large.

- The **numerator** in the partial F test measures the extra variation in the data not explained by the reduced model, but explained by the full model.
- The **denominator** is the variation in the data not explained by the full model (i.e., not explained by either model), which is used to estimate the error variance.

Birthweight Example (cont'd)

We wish to test whether the group of predictors that describe the father's characteristics (i.e. variables $X_9 - X_{12}$) is jointly significant. The hypotheses are

$$\begin{cases} H_0 : Y \sim 1 + \beta_1 X_1 + \dots + \beta_8 X_8 \\ H_\alpha : Y \sim 1 + \beta_1 X_1 + \dots + \beta_8 X_8 + \dots + \beta_{12} X_{12} \end{cases}$$

Here $p_0 = 8 + 1$ (8 predictors plus intercept) and $p_\alpha = 12 + 1$ is the number of β 's in the full model. Extracting RSS_0 and RSS_α in R/Python and computing the partial F test statistic yields $F = 0.7225$. Because this value is smaller than the critical value $F_{4, n-13} = 2.7$ we do not reject H_0 and prefer the reduced model.

Partial F Test: Special Cases

- **Test for a Single Predictor β_j**

$$\begin{cases} H_0 : Y \sim 1 + \beta_1 X_1 + \dots + \beta_{j-1} X_{j-1} + \beta_{j+1} X_{j+1} + \dots + \beta_p X_p \\ H_\alpha : Y \sim 1 + \beta_1 X_1 + \dots + \beta_{j-1} X_{j-1} + \beta_j X_j + \beta_{j+1} X_{j+1} + \dots + \beta_p X_p \end{cases}$$

The partial F test for a single coefficient β_j is exactly equivalent to the usual t -test

$$t_j = \frac{\hat{\beta}_j}{\widehat{\text{se}}(\hat{\beta}_j)} \sim t_{n-p-1}$$

under $H_0 : \beta_j = 0$. In practice most software reports the t -statistic and its p -value directly.

- **Test for all Predictors**

$$\begin{cases} H_0 : Y \sim 1 & [\text{intercept-only model}] \\ H_a : Y \sim 1 + \beta_1 X_1 + \dots + \beta_{j-1} X_{j-1} + \beta_j X_j + \beta_{j+1} X_{j+1} + \dots + \beta_p X_p \end{cases}$$

This is the **overall** F test, which can be viewed as a *goodness-of-fit test* for the regression model.

2.8 Categorical Variables in MLR

As we discussed, some of the predictors may be *categorical*. Consider a categorical predictor with k levels. When it is included in a linear regression model it is represented by binary indicator variables. Two common coding schemes are used:

One-hot encoding creates k binary indicators, one for each level:

$$D_i = \begin{cases} 1 & \text{if the observation belongs to level } i, \\ 0 & \text{otherwise.} \end{cases}$$

Dummy coding (the default in classical linear regression) creates only $k-1$ binary indicators and absorbs one level into the intercept. The indicators take the form

$$D_i = \begin{cases} 0 & \text{if the observation is not at level } i, \\ 1 & \text{if the observation is at level } i, \end{cases}$$

where the omitted level is called the reference level.

An Example with Categorical Variables

Consider a categorical predictor with three levels (a , b , c).

One-hot encoding produces $k = 3$ columns:

$$\begin{pmatrix} a \\ a \\ b \\ b \\ b \\ c \\ c \end{pmatrix} \Rightarrow \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{pmatrix}$$

Dummy coding (with level c as the reference) produces $k - 1 = 2$ columns:

$$\begin{pmatrix} a \\ a \\ b \\ b \\ b \\ c \\ c \end{pmatrix} \Rightarrow \begin{pmatrix} 1 & 0 \\ 1 & 0 \\ 0 & 1 \\ 0 & 1 \\ 0 & 1 \\ 0 & 0 \\ 0 & 0 \end{pmatrix}$$

In the dummy-coding matrix, Column 1 is the indicator for level a and Column 2 is the indicator for level b . Level c is absorbed into the intercept.

In general any level may be chosen as the reference level; the choice is usually driven by subject-matter context rather than by statistical considerations.

Dummy coding vs. One-hot encoding

- Dummy coding uses $k - 1$ indicators and keeps the design matrix full rank when an intercept is present. This is the standard approach in classical linear regression.
- One-hot encoding uses k indicators. When an intercept is also included, the k columns sum to the constant vector $\mathbf{1}$, creating perfect multicollinearity; $\mathbf{X}^T \mathbf{X}$ becomes singular and the ordinary least-squares estimator is no longer unique.
- Fitted and predicted values are identical under both codings. *Only the interpretation of the individual coefficients changes.* In practice, any level may be chosen as the reference; the choice is usually guided by subject-matter context rather than by statistical considerations.

2.9 Collinearity

In practice we often encounter data sets in which many of the predictors are highly correlated with one another. In such cases the fitted values and the sampling variances of the regression coefficients can depend heavily on which particular predictors are included in the model. The resulting estimators become sensitive and unreliable.

Possible *symptoms* of collinearity include :

- (i) high pairwise (sample) correlation among the predictors,
- (ii) a relatively large R^2 ,

- (iii) a statistically significant overall F test even though non of the individual predictors is statistically significant.

Specific diagnostic tests and metrics exist for quantifying the severity of collinearity, but they lie beyond the scope of this course.

The simplest remedy is to remove one or more predictors from each highly correlated group. More advanced approaches include principal-component analysis (PCA) and regularization methods based on penalized least squares.

2.10 Model Diagnostics

We made several assumptions when defining the linear regression model

$$\mathbf{y} = \mathbf{X}\beta + \varepsilon, \text{ where } \varepsilon \sim IID(0, \sigma^2)$$

These assumptions are:

- *Linearity*: the conditional mean of Y is a linear function of the predictors.
- *Constant Variance*(homoscedasticity): all responses share a common, constant variance σ^2 .
- *Uncorrelated errors*: the errors ε_i and ε_j (and therefore Y_i and Y_j) are uncorrelated for $i \neq j$.
- *(Normality)*: this is not required for the derivation of the least-squares estimators themselves. It becomes necessary, however, when we want exact finite-sample distributional results for the estimators, fitted values, and residuals, and it is especially important for hypothesis testing.

If the goal is to interpret the effects of the predictors on the response (rather than pure prediction), all of the above assumptions should be checked with diagnostic plots and tests, and appropriate remedial measures should be taken when they are violated.

Because the primary focus of this course is prediction, we refer students to a dedicated regression course or textbook for a thorough treatment of diagnostics and remedial measures in multiple linear regression.

Outliers

Outliers can be a concern in both large and small data sets. When the sample size is modest, it is advisable to test formally for outliers—for example, with tests based on the leave-one-out prediction error. In very large data sets such

tests lose power and are less useful; one should instead adjust for multiple comparisons or adopt alternative strategies.

If outliers appear to be problematic, useful practical steps include:

- (i) examining the range of each variable,
- (ii) applying a log, square-root, or other suitable transformation to right-skewed predictors or to the response Y , and
- (iii) using winsorization to limit the influence of extreme values.

2.11 The Birthweight Data Set Example

The `birthweight.csv` data set contains observations on 42 babies. The variables are:

- **ID**: Unique Identification number of a baby
- **Length**: Length of the baby at time of birth in cm (X_1)
- **Birthweight**: Weight of the baby at time of birth in kg (Y)
- **Headcirc**: Head circumference of the baby at time of birth in cm (X_2)
- **Gestation**: Completed weeks of gestation (X_3)
- **smoker**: Mother is/is not a smoker {0: No, 1: Yes} (X_4)
- **mage**: Mother's age at time of birth (X_5)
- **mnocig**: Mother's number of cigarettes smoked per day (X_6)
- **mheight**: Mother's height in cm (X_7)
- **mppwt**: Mother's pre-pregnancy weight (X_8)
- **fage**: Father's age at time of birth (X_9)
- **fedyrs**: Father's years of education (X_{10})
- **fnocig**: Father's number of cigarettes smoked per day (X_{11})
- **fheight**: Father's height (X_{12})

You can download the data file here: [Birthweight.csv](#) .

The analysis of this data set is available in R and Python.

The data analysis example also includes a highly correlated data set (`seatpos`) from Faraway's book on "Linear Models in R". You can download this data here: [seatpos.csv](#)